

ZipperGen Workflow Development and Deployment Manual

Concepts, commands, configuration, runtime, and operations

The ZipperGen Authors

Manual edition: 8 August 2026

Abstract

This manual is what you need after the tutorial. It covers the project's files, the workflow language, inspection and validation, running and testing, models and connectors, deployment and operations, and what to do when something is wrong.

It assumes you can use a terminal and an editor. It does not assume you are a Python or operations engineer, and it does not assume you use a coding agent — though it says where one helps.

Start with the tutorial

`docs/first-workflow.pdf` builds one small application end to end in about seven pages. This manual explains the parts it passes over. If you have not read it, read it first; most of what follows will then be lookup rather than learning.

The promise

The workflow is always ordinary Python. A coding agent may write or change it, but ZipperGen does the deterministic work: loading it, projecting it for every participant, checking that the result is well formed, running it durably, and deploying it. Opaque model output never becomes a deployed application without passing through visible code you can read.

Contents

1	The project	3
1.1	<code>zippergen.toml</code>	3
1.2	What stays on one machine	4
2	Writing workflows	4
2.1	Participants and messages	4
2.2	Actions	4
2.3	Decisions	5
2.4	Parallel regions and co-regions	5
2.5	Inputs and outputs	5

3	Inspecting and validating	5
3.1	<code>zg validate</code>	5
3.2	<code>zg show</code>	6
4	Running and testing	6
4.1	One verb	6
4.2	Choosing a model	7
4.3	Deterministic testing	7
4.4	Durable runs	7
5	Models, people, and services	8
5.1	Models	8
5.2	Asking a person	8
5.3	Reading and writing external services	8
5.4	Authorizing Google	9
6	Deployment	9
6.1	Prepare, then start	9
6.2	Operating one	10
6.3	What removal keeps	10
6.4	Moving a project to a server	10
7	Working with a coding agent	10
8	When something is wrong	11
8.1	The workflow will not load	11
8.2	A participant seems to be waiting forever	11
8.3	Only one branch ever runs	11
8.4	A deployment refuses to start	11
8.5	<code>zg status</code> shows nothing after a run	11
8.6	Connector configured but the question still appears in the terminal	12
9	Command reference	12
10	Where the guarantee comes from	12

1 The project

A ZipperGen project is a directory. `zg init` creates one:

File	In git	Holds
<code>specification.md</code>	yes	what the workflow should do, in prose
<code>workflow.py</code>	yes	the coordination, as Python
<code>zippergen.toml</code>	yes	models, connectors, entry point — no secrets
<code>AGENTS.md</code>	yes	points a coding agent at <code>zg skill</code>
<code>.zippergen/</code>	no	scratch
<code>ZIPPERGEN_HOME</code>	no	credentials, runs, deployments

The split that matters is the last two rows. Everything above them travels with the project; everything in them describes one machine. A colleague who clones the repository gets the workflow, the specification and the configuration, and supplies their own keys.

1.1 zippergen.toml

```

1  schema_version = 1
2  name = "email-approval"
3  specification_file = "specification.md"
4  workflow_entry = "workflow.py:email_approval"
5
6  [models.configurations."openai-gpt-4o"]
7  provider = "openai"
8  model = "gpt-4o"
9  spec = "openai:gpt-4o"
10
11 [models.assignments]
12 default = "mock"
13
14 [models.assignments.lifelines]
15 Writer = "openai-gpt-4o"
16
17 [connectors.configurations."approvals"]
18 provider = "telegram"
19 kind = "telegram"
20 chat_id = "12345678"
21
22 [connectors.assignments.lifelines]
23 User = "approvals"
```

It is an ordinary file. Edit it in an editor, or let a coding agent edit it; nothing in ZipperGen owns it. `workflow_entry` is why most commands do not need you to name the workflow.

Named configurations exist so several participants can share one and it changes in one place. `configurations` says what a thing is; `assignments` says who uses it.

1.2 What stays on one machine

API keys	environment variables, or ZIPPERGEN_HOME
Telegram bot token	typed once, never an argument
Google credentials	<code>zg connector authorize / accept</code>
Local model endpoints	site configuration
Runs and deployments	ZIPPERGEN_HOME

2 Writing workflows

A workflow is a Python function decorated with `@workflow`. Its body is read, rewritten, and turned into an intermediate representation — so it looks like Python and reads like a protocol, but it is not executed line by line as written.

2.1 Participants and messages

```

1 User = Lifeline("User")
2 Writer = Lifeline("Writer")
3
4 User(message) >> Writer(message)
```

The left side sends, the right receives. The names in brackets say which variables carry the value at each end; they need not match.

A self-send (`A(x) >> A(y)`) is a local rename, not a message.

2.2 Actions

An action is work one participant does.

<code>@llm</code>	a model call with a prompt and declared outputs
<code>@pure</code>	ordinary Python; output name and type from the return annotation
<code>@human</code>	stops and asks a person: <code>confirm</code> , <code>select</code> , <code>input</code>
<code>@assistant</code>	runs a coding agent against a repository
<code>@planner</code>	asks a model to generate a sub-workflow, validated before it runs

```

1 @llm(
2     system="You write short, friendly replies.",
3     user="Reply to this message:\n\n{message}",
4     parse="text",
5     outputs=[("draft", str)],
6 )
7 def draft_reply(message: str): ...
8
9 Writer: draft = draft_reply(message)
```

Several actions by one participant can be grouped:

```

1 with LLM1:
2     agreed = check_agreement(verdict, other_verdict)
3     trials = inc_trials(trials)
```

2.3 Decisions

```

1 if approved @ User:
2     User: result = sent(draft)
3 else:
4     User: result = discarded()

```

The `@ User` names the participant who evaluates the condition. It is required, and it is the most important annotation in the language: it decides who broadcasts the outcome and who merely receives it.

Loops work the same way:

```

1 while (not agreed and trials < MAX_ROUNDS) @ LLM1:
2     ...

```

Parenthesise compound conditions

`@` binds more tightly than `not`, `and` and `or`. Write `if (not agreed) @ LLM1:`, not `if not agreed @ LLM1:`.

2.4 Parallel regions and co-regions

Independent branches that may interleave:

```

1 with parallel:
2     with branch:
3         Orchestrator(candidate) >> TestRunner(candidate)
4         TestRunner: (test_status,) = run_tests(candidate)
5         TestRunner(test_status) >> Committer(test_status)
6     with branch:
7         Orchestrator(candidate) >> Security(candidate)
8         Security: (security_status,) = scan_security(candidate)
9         Security(security_status) >> Committer(security_status)

```

A co-region is weaker: messages that may arrive in any order.

```

1 with coregion:
2     Analyst_A(a) >> Collector(a_report)
3     Analyst_B(b) >> Collector(b_report)

```

2.5 Inputs and outputs

```

1 def email_approval(message: str @ User) -> str:
2     ...
3     return result @ User

```

An input belongs to the participant that supplies it; the return says who holds the result.

3 Inspecting and validating

3.1 zg validate

```
zg validate
```

Loads the module, projects the protocol for every participant, checks the deployment declaration, and renders the result back. If it passes, the workflow is well formed and the deadlock-freedom guarantee applies to it.

This is the check to run after every change. It is fast and needs nothing configured.

3.2 zg show

<code>zg show</code>	the global protocol
<code>zg show --communications</code>	messages only, no action bodies
<code>zg show --agent NAME</code>	one participant's local program
<code>zg show --detail full</code>	everything, including action definitions
<code>zg show --format json</code>	structured, for tooling

`--agent` is the interesting one. It shows what a participant actually runs after projection, which is not what you wrote:

```
1 @role('Writer')
2 def email_approval__Writer() -> None:
3     message = recv('User')
4     draft = draft_reply(message)
5     send('User', draft)
```

The `Writer` has no branch, because it takes no part in the `User`'s decision. A participant's local program mentions only what concerns it — which is why it cannot wait for a message the decider chose not to send.

Use this when a workflow behaves in a way you did not expect. The global view says what you wrote; the local view says what will run.

4 Running and testing

4.1 One verb

<code>zg run</code>	execute once; nothing is kept
<code>zg run --durable</code>	record every step, so it can be resumed
<code>zg run --resume</code>	continue the most recent unfinished run
<code>zg run --resume --run-id ID</code>	continue a particular one

A plain run leaves nothing behind, so `zg status` afterwards finds nothing. That is deliberate: throwaway runs should not accumulate. When you want to come back to a run — to resume it, inspect it, or answer a human action later — use `--durable`.

Missing inputs are asked for in a terminal. In a script, supply them:

```
zg run --input message="Could we move our meeting to Thursday?" --yes
```

4.2 Choosing a model

<code>--llm mock</code>	a placeholder for every action; no key needed
<code>--llm openai:gpt-4o-mini</code>	a real provider
<code>--llm scripted:replies.json</code>	recorded answers; see below
<code>--llm-for Writer=openai:gpt-4o</code>	override one participant or action

Without `--llm`, the assignments in `zippergen.toml` apply.

4.3 Deterministic testing

`mock` answers every action the same way, so a workflow driven by it takes one path and no other. In a consensus protocol that means immediate agreement; everything behind a decision is unreachable.

Scripted responses fix that:

```

1 {
2   "LLM1.assess":    {"verdict": "yes", "reason": "unmoved"},
3   "LLM2.assess":    {"verdict": "no",  "reason": "unmoved"},
4   "LLM1.reconsider": {"verdict": "yes", "reason": "unmoved"},
5   "LLM2.reconsider": {"verdict": "no",  "reason": "unmoved"}
6 }
```

```
zg run --llm scripted:never-agree.json
```

Keys are `action` or `Participant.action`; the more specific one wins. Values follow one rule worth learning:

<code>{...}</code>	a constant: every call is answered the same way
<code>[{...}, {...}]</code>	a finite sequence: exactly these calls are expected

Running past the end of a list is an error, not a silent repeat. That is the point: if a change makes an action run more often than the script expects, the run fails instead of passing quietly.

Use it for branches, not for prompts

Scripted responses exercise *coordination*: this reviewer disagrees, that loop runs three times, this branch is taken. They say nothing about whether a real model would produce good text. Both matter; they are different questions.

4.4 Durable runs

A durable run records each step before taking it, in its own SQLite store. An interrupted run resumes where it stopped, and a model call already made is not made again.

```

zg run --durable --llm openai:gpt-4o-mini
# Ctrl-C
zg run --resume
```

running	executing now
waiting	stopped at a human action
interrupted	you pressed Ctrl-C
failed	a participant raised an error
done	finished, with a result

interrupted and failed can be resumed; done cannot.

5 Models, people, and services

5.1 Models

For development, an environment variable is enough:

```
export OPENAI_API_KEY=sk-...
zg run --llm openai:gpt-4o-mini
```

openai:MODEL	OPENAI_API_KEY
anthropic:MODEL (or claude:)	ANTHROPIC_API_KEY
mistral:MODEL	MISTRAL_API_KEY
ollama:MODEL (or local:)	OLLAMA_BASE_URL
mock	—

To fix the choice in the project rather than the command line, put it in `zippergen.toml` under `models`.

5.2 Asking a person

A `@human` action stops and waits. Where the question appears depends on configuration:

nothing configured	the terminal running the workflow
a Telegram connector, assigned	a message with buttons

The terminal is right while developing and wrong for a deployment nobody is watching. Two commands, doing two different things:

```
zg connector configure telegram approvals --chat-id 12345678
zg connector assign User approvals
```

The first says *which* chat. The second says *who* is asked there — a participant, or one action with `User.approve_reply`.

The bot token is typed without echo and stays on this machine. The chat id goes in `zippergen.toml` and is committed: a colleague gets the chat, and supplies their own token.

5.3 Reading and writing external services

A workflow declares what it needs:


```

1 zippergen_connectors = (
2     ConnectorRequirement(
3         name="call-mailbox", kind="gmail",
4         participant="Mailbox", access="read-write",
5     ),
6 )

```

and the project says which one, binding it to that requirement by name:

```

zg connector configure gmail inbox \
    --query "is:unread in:inbox" --bind call-mailbox
zg connector configure google-sheets records \
    --spreadsheet-id 1AbC... --tab Calls --bind call-records

```

5.4 Authorizing Google

```
zg connector authorize google --scopes gmail.modify,spreadsheets
```

This opens a browser, and prints an encoded result rather than saving it. On the machine that will run the workflow:

```
zg connector accept google
```

which asks for the result without echo and stores the credential together with the scopes Google actually granted.

The two halves exist because a server usually has no browser. Authorize on your own computer; accept on the server. The credential never appears in a command line, so it does not reach shell history.

Scopes are checked before deployment

Deploying refuses when the granted scopes do not cover what the workflow needs — for example a workflow that marks mail as read needs `gmail.modify`, not `gmail.readonly`. Re-run `authorize` with the missing scopes.

6 Deployment

6.1 Prepare, then start

```
zg deploy --name production --no-start
zg start production
```

These are deliberately separate. Preparing writes the bundle, a managed Python environment, and the service files, and reaches nothing — a deployment can be prepared today and started next week. Starting is about to make real calls, so every model and connector is probed first and a failure stops it.

After the first time, the deployment has a name and you use that:

```
zg deploy production          # redeploy the current source
```

6.2 Operating one

<code>zg status NAME</code>	is it running, and where is it
<code>zg logs NAME</code>	what it has been doing
<code>zg trace NAME</code>	recent protocol events
<code>zg tasks NAME</code>	decisions waiting for a person
<code>zg approve NAME</code>	answer one
<code>zg doctor NAME</code>	readiness checks, in detail
<code>zg stop / restart NAME</code>	move the service
<code>zg compact NAME</code>	reclaim space; the service must be stopped
<code>zg remove NAME</code>	delete it

6.3 What removal keeps

durable store, log	kept	the record of what actually ran
profile	kept	what produced them
secrets	deleted	must not be left behind
environment, bundle	deleted	rebuilt by deploying again
service files	deleted	regenerated

`--purge` keeps nothing, including the store. What survives lands in `ZIPPERGEN_HOME/trash/deployments/`, readable only by you. Nothing prunes it, so look occasionally.

6.4 Moving a project to a server

1. Clone the repository. The workflow, specification and `zippergen.toml` come with it.
2. `zg validate` — confirms the module loads in that environment.
3. Supply what the machine needs: model keys in the environment, the Telegram token via `connector configure`, Google via `authorize` on your laptop and `accept` there.
4. `zg deploy --name production --no-start`, then `zg start production`.

Nothing secret is in the repository, so step 3 is the whole difference between a clone and a running service.

7 Working with a coding agent

ZipperGen ships instructions for coding agents. `zg init` writes an `AGENTS.md` pointing at them; `zg skill` prints them.

```
zg skill          # what an agent should follow
zg skill --agents-md  # a pointer file for an existing project
```

A reasonable division:

the agent	writes and changes the workflow and specification, validates, runs, inspects projections, prepares deployments
you	credentials, authorization, and starting production

That division is a convention, not a wall: an agent with a shell can run any command you can, and can read any file you can. Credentials are typed rather than passed as arguments, so they stay out of the agent’s conversation, your shell history and the repository — but once saved they live in a file under `ZIPPERGEN_HOME` that your user can read.

That distinction is worth keeping straight. Entering a credential interactively protects it *at the moment of entry*; it is not a capability boundary. For a real one, run the agent under a different account, or in an environment that does not carry `ZIPPERGEN_HOME`. Nothing in ZipperGen can supply that, because a process running as you can do what you can do.

The specification is yours and the agent’s

ZipperGen does not check that `workflow.py` implements `specification.md`, because it cannot: one is prose. There is no gate, no provenance, and no ceremony around it. Keep it accurate the way you keep a README accurate — and tell your agent to update it before changing behaviour, which the shipped skill does.

8 When something is wrong

8.1 The workflow will not load

`zg validate` reports the exception. Common causes: a missing import, an action used before it is defined, or a compound condition without parentheses (`if (not agreed) @ LLM1:`).

8.2 A participant seems to be waiting forever

Look at what it actually runs:

```
zg show --agent ThatParticipant
```

A projected program cannot wait for a message nobody sends, so if a participant is stuck the cause is usually an action that never returns — a model call without a timeout, or a human action nobody has answered. `zg tasks` lists the second kind.

8.3 Only one branch ever runs

You are almost certainly using `--llm mock`, which answers every action identically. Script the responses instead.

8.4 A deployment refuses to start

Starting probes every model and connector. The message names what failed. The usual causes are a missing API key on that machine, a Google authorization that does not cover the workflow’s scopes, or a Telegram token that was never supplied there.

8.5 `zg status` shows nothing after a run

A plain run keeps nothing. Use `--durable` for a run you want to come back to.

8.6 Connector configured but the question still appears in the terminal

Configuring says which chat; assigning says who is asked there. Both are needed:

```
zg connector assign User approvals
```

9 Command reference

Project	
init [NAME]	create a project here
skill	print the coding-agent instructions
Inspect	
validate	load, project, and check
show	render the protocol or one participant
snapshot / diff	record and compare semantics
Run	
run	execute; <code>--durable</code> , <code>--resume</code>
status / trace / tasks / approve	inspect and answer
Configure	
connector configure telegram gmail google-sheets	save one
connector assign TARGET NAME	route human actions
connector authorize google	browser flow, prints a result
connector accept google	save that result here
Deploy	
deploy [--name NAME] [--no-start]	prepare, and start
start / stop / restart	move the service
logs / doctor	operate
compact / remove	reclaim, delete

Every command takes `--help`. Commands that act on a workflow default to the project's; commands that act on a deployment default to the most recent.

zippergen and zg are the same program.

10 Where the guarantee comes from

Projection is syntax-directed. For each decision, every participant is one of three things:

decider	evaluates the condition and broadcasts the outcome
recipient	appears in a branch; receives the outcome and follows it
bystander	appears in neither branch; the decision is erased from its program

The **Writer** in the tutorial is a bystander. Its local program does not mention the approval, so it cannot wait on it or disagree about it.

The main theorem states that the projected local programs produce exactly the behaviours of the global workflow, up to the control messages projection introduces. Deadlock freedom follows for well-formed workflows. Both are machine-checked in Lean 4.

That is the trade the whole system makes: write one protocol, accept the language’s restrictions, and get a distributed implementation that cannot deadlock — rather than writing each agent separately and hoping.